

Estimating statistical model parameters with maximum likelihood in R

As was covered in lecture there is a general recipe for estimating statistical model parameters in Python.

1. load data
2. write a custom likelihood function
3. estimate parameter values by minimizing the negative log likelihood


The code required to accomplish these steps on a simulated data set are below.

```
In [1]: ### Import packages
import numpy
import pandas
from scipy.optimize import minimize
from scipy.stats import norm
from plotnine import *

### Simulating data
# creating a uniformly distributed set of values for an independent variable x
# and values for a variable y that is linearly dependent on x
N=25
x=numpy.random.uniform(0,20,size=N)
y=3*x+5

# add some "noise" to y and put the variables in a dataframe
y=y+numpy.random.randn(N)*3
df=pandas.DataFrame({'x':x,'y':y})

# plot our observations
ggplot(df,aes(x='x',y='y'))+geom_point()+theme_classic()
```



Out[1]: <ggplot: (297143721)>

In [2]: *### Custom likelihood function*

```
def nllike(p,obs):
```

```
    B0=p[0]
```

```
    B1=p[1]
```

```
    sigma=p[2]
```

```
    expected=B0+B1*obs.x
```

```
    nll=-1*norm(expected,sigma).logpdf(obs.y).sum()
```

```
    return nll
```

```
### estimate parameters by minimizing the negative log likelihood
```

```
initialGuess=numpy.array([1,1,1])
```

```
fit=minimize(nllike,initialGuess,method="Nelder-Mead",options={'disp': True},args=df)
```

```
# fit is a variable that contains an OptimizeResult object
```

```
# attribute 'x' is a list of the most likely parameter values
```

```
print(fit.x)
```

Optimization terminated successfully.

Current function value: 65.715476

Iterations: 133

Function evaluations: 231

```
[ 5.47260034  3.04434477  3.35240437]
```